

# COMP3420 Lesson 13

Greg Baker

2023-05-29



# Today's half-session

- A checklist of learning outcomes for the course
- Walk-through of a sample exam

# Week 7

# Memorize

Things you should **memorize** (because you will use them in real life, too):

- Regular expressions before PCRE (i.e. before Perl)
  - \* [] ? . ^ \$ [^A-Z] [a-z] abc \
- The names of code pages for languages that you speak:
  - e.g. CP1252, ISO-8859-1 (Western European languages)

In 5 years' time, email me to say thanks when you realise you are using these things all the time.

# Be able to define

The golden rule of text *When you have a stream of bytes, and you don't know how it was encoded, you have a useless stream of bytes.*

Mojibake What it is (and what it looks like)

æ-†â—â€-ã '   
 â€œâ€œ...; â/f-â/f\*â/f14q™%ç§'â"â...â€Žâ;â,çâ,â/fââf†â,çâ,eï%Wikipedia!%â€

Fixed-width encoding vs varying width encoding

# Python code you should be able to explain

`re.compile()` Regex compile

`x.decode('ascii')` Take a binary and interpret it as ASCII

`x.encode('utf-8')` Take a string and make it UTF-8

`open('filename', encoding='utf-8')` Read a file with an encoding  
except `UnicodeDecodeError`: Even when you aren't using Unicode.

# Terms to know: Unicode Transformation Formats

- UTF-32** Fixed length encoding, 4 bytes for every character (space inefficient); just store the Unicode code point.
- UTF-16** Varying length encoding, mostly used in Microsoft Windows and JavaScript. Incompatible with ASCII.
- UTF-8** Varying length encoding, compatible with ASCII, most widely used.

# Manipulate: UTF-8 $\Leftrightarrow$ Unicode points

(Don't need to memorise this, just understand how it works)

First code point	Last code point	Byte 1	Byte 2	Byte 3	Byte 4
U+0000	U+007F	0xxxxxxx			
U+0080	U+07FF	110xxxxx	10xxxxxx		
U+0800	U+FFFF	1110xxxx	10xxxxxx	10xxxxxx	
U+10000	U+10FFFF	11110xxx	10xxxxxx	10xxxxxx	10xxxxxx



# Week 8

# Define

**Vocabulary** Number of distinct words in a corpus

**Corpus size** Number of words in a corpus

**Happax legomonon** A word that appears only once in a corpus —  
typically 40%-60% of words

**Bigram** Two word sequences

**Stop word** Words that get ignored

# Analyse and Interpret Parameters

Zipf's Law  $f(r) = \frac{C}{r^s}$

Heap's Law, Herdan's Law  $V = kN^\beta$

Search term hit rate estimation  $\frac{N}{V} = \frac{N}{kN^\beta} = \frac{N^{1-\beta}}{k}$

Term usefulness estimate for a classifier  $\frac{C}{V} = \frac{N}{kN^\beta L} = \frac{N^{1-\beta}}{kL}$

(The last two are not in the class textbooks.)

- No need to memorise these; in real life you can look them up!
- Be able to explain the kinds of corpora that would lead to different values for  $s$  and  $\beta$ , and the implications for search, classification and author identification.

# Python code you should be able to explain

`nltk.sent_tokenize(x)` Break document into sentences.

`nltk.word_tokenize` Break sentence or document into words using rules typical of English-language texts.

`tfidf = TfidfVectorizer(); vecs = tfidf.fit_transform(x)` sklearn's vectorizer, creating TFIDF vectors

`sklearn.metrics.pairwise.cosine_similarity(x,y)` A function to measure similarity between two vectors

# Manipulate

Be able to perform one iteration of Byte-pair encoding by hand.  
(Not really a useful skill, but if you can do it, you will have an intuition for how and why it works.)

Repeat:

- Choose the two symbols that are most frequently adjacent in the training corpus (say 'A', 'B')
- Add a new merged symbol 'AB' to the vocabulary
- Replace every adjacent 'A' 'B' in the corpus with 'AB'.

Until  $k$  merges have been done, or the vocabulary is the target size.

# Week 9

# Python code you should be able to explain

`Dense(1, activation='sigmoid')` Also known as logistic regression.

`Dense(10, activation='softmax')` The last layer in a multi-class classification problem.

`Dense(20, activation='relu')` A layer which creates regions out of a dataset.

`t = TextVectorization(); t.adapt(x)` The Keras way of turning documents into vectors

# Analyse and interpret

- `vocab = vectorizer.get_vocabulary()`
- `weights = model.get_weights()[0][:,0]`
- `print(zip(vocab,weights))`

Weights in a logistic regression describe the impact of a word in a document towards one classification or the other. (Big positive numbers = “strongly associated with that class”)



## Be able to define

**Data irrelevancy** Nothing in  $X$  predicts  $y$ .

**True positive, true negative, false positive, false negative** What the system predicted vs what ground truth said

**Precision, recall, accuracy** Metrics for evaluating a classifier

**Overfitting** The model is memorizing the data rather than generalizing

**Underfitting** The model is unable to get good results because it isn't capture

# Terms to Know

**GDPR and PIPL** EU and Chinese laws that have *extra-territoriality* and require explainability (even if your company is in Australia) if the model is used for something important that might affect someone's life.

**F1 score** A balanced score that captures both precision and recall. Trading off precision for recall will generally worsen F1 score

## Be able to give examples of

- A classifier that would need explainability under GDPR or PIPL.
- Reasons you might trade off explainability vs accuracy
- An unbalanced data set
- Data sets that you would use each metric for
- Things you might do to reduce overfitting and underfitting

# Week 10

# Define

- Embedding** A mapping between language (usually words) and numeric vectors
- Contextual embedding** An embedding that distinguishes between different meanings for the same written word
- Non-contextual embedding** An embedding that uses the written form and doesn't distinguish meaning
- Bag-of-words** Vectorise by giving each word its own dimension
- Hypernym / hyponym** A hypernym is a more general concept; a hyponym is a specific example. "Animal" is a hypernym of "dog".
- Null region** The volume of linear space where all Relu functions are zero, and all points are indistinguishable.
- Context drift** Models perform worse over time; the future brings new words and the task may change

## Python code you should be able to explain

- `s = wordnet.synsets(x)` What different meanings does the word `x` have?
- `s[0].definition()` For the first meaning, what is its definition?
- `s[0].lemmas()` What words express the first meaning?
- `s[0].hypernyms()[0]` What is the first hypernym of the first meaning?

# Be able to give examples of

- Problems with the bag-of-words approach
- Ways to map words into numeric vectors, and whether they are
- Context and non-contextual embeddings

## Python code you should be able to explain

```
c=Constant(embedding_matrix)]  
embedding_layer = keras.layers.Embedding(  
    input_dim=len(voc),  
    output_dim=50,  
    embeddings_initializer=c,  
    trainable=False)
```

What does trainable=False mean? Why would we set this?



# Week 11

# Be able to define

**Temperature** A controllable parameter used in a sequence-to-sequence which increases the randomness of the selection from the probability distribution of next tokens.

**Sequence-to-sequence model / text generation** A model designed to predict the next character, word or token in a stream

**AI alignment problem** the challenge of designing artificial intelligence systems that reliably understand and execute human intentions and values, even when operating at superhuman levels of performance.

# Misc

- Recognize architecture diagrams for RNNs, LSTM and Transformers
- Predict the results of changing the temperature of a model's output
- Follow Chollet's Guidance to identify the most likely-to-be-effective architecture for a classifier problem. i.e. Make a choice of classifier given the following guidance:  

$$\frac{\text{Number of samples}}{\text{Mean sample length}}$$
  - < 1500 Bag of words + logistic regression
  - < 15,000 Pretrained embeddings + LSTM or Transformer
  - > 15,000 Learn embeddings + Transformer

# Week 12

## Give examples of

- Differences between GPT-3.5 and GPT-4
- Prompt injection
- The purpose and effects of RLHF (Reinforcement learning from human feedback)
- Tasks that GPT-4 is unable to do

# Define

**Prompt injection** If a prompt to a language model includes text from a third party, that third party can manipulate the language model to produce uncontrollable responses.

**Context length** The amount of input text the model can take into account when generating a response or prediction.

**Hallucination** The model generates output that is plausible but untrue.